

▼ We want to predict the price of houses by creating a model that will help us predict the prices

▼ Let's start by importing the libraries needed for the project

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set_style("whitegrid")
import warnings
warnings.filterwarnings("ignore")
```

▼ Let's load the datasets

```
"""df1 = pd.read_csv("test.csv")"""
```

```
df = pd.read_csv("train.csv")
```

```
df.shape
```

```
(1460, 81)
```

▼ EDA

```
df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	Land
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

5 rows × 81 columns



▼ Missing Values

```
df.isnull().sum()
```

```
Id          0
MSSubClass  0
MSZoning    0
LotFrontage 259
LotArea     0
...
MoSold      0
YrSold      0
SaleType    0
SaleCondition 0
SalePrice   0
Length: 81, dtype: int64
```

▼ Taking care of the missing values

```
df.dropna(axis=1, inplace = True)
```

```
df.shape
```

```
(1460, 62)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 62 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id              1460 non-null  int64
1   MSSubClass      1460 non-null  int64
2   MSZoning        1460 non-null  object
3   LotArea         1460 non-null  int64
4   Street          1460 non-null  object
5   LotShape        1460 non-null  object
6   LandContour     1460 non-null  object
7   Utilities       1460 non-null  object
8   LotConfig       1460 non-null  object
9   LandSlope       1460 non-null  object
10  Neighborhood    1460 non-null  object
11  Condition1      1460 non-null  object
12  Condition2      1460 non-null  object
13  BldgType        1460 non-null  object
14  HouseStyle      1460 non-null  object
15  OverallQual     1460 non-null  int64
16  OverallCond     1460 non-null  int64
17  YearBuilt       1460 non-null  int64
```

18	YearRemodAdd	1460	non-null	int64
19	RoofStyle	1460	non-null	object
20	RoofMatl	1460	non-null	object
21	Exterior1st	1460	non-null	object
22	Exterior2nd	1460	non-null	object
23	ExterQual	1460	non-null	object
24	ExterCond	1460	non-null	object
25	Foundation	1460	non-null	object
26	BsmtFinSF1	1460	non-null	int64
27	BsmtFinSF2	1460	non-null	int64
28	BsmtUnfSF	1460	non-null	int64
29	TotalBsmtSF	1460	non-null	int64
30	Heating	1460	non-null	object
31	HeatingQC	1460	non-null	object
32	CentralAir	1460	non-null	object
33	1stFlrSF	1460	non-null	int64
34	2ndFlrSF	1460	non-null	int64
35	LowQualFinSF	1460	non-null	int64
36	GrLivArea	1460	non-null	int64
37	BsmtFullBath	1460	non-null	int64
38	BsmtHalfBath	1460	non-null	int64
39	FullBath	1460	non-null	int64
40	HalfBath	1460	non-null	int64
41	BedroomAbvGr	1460	non-null	int64
42	KitchenAbvGr	1460	non-null	int64
43	KitchenQual	1460	non-null	object
44	TotRmsAbvGrd	1460	non-null	int64
45	Functional	1460	non-null	object
46	Fireplaces	1460	non-null	int64
47	GarageCars	1460	non-null	int64
48	GarageArea	1460	non-null	int64
49	PavedDrive	1460	non-null	object
50	WoodDeckSF	1460	non-null	int64
51	OpenPorchSF	1460	non-null	int64
52	EnclosedPorch	1460	non-null	int64
53	3SsnPorch	1460	non-null	int64

```
df.describe()
```

	Id	MSSubClass	LotArea	OverallQual	OverallCond	YearBuilt
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000

```
df.columns
```

```
Index(['Id', 'MSSubClass', 'MSZoning', 'LotArea', 'Street', 'LotShape',
      'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood',
      'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'OverallQual',
      'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle', 'RoofMatl',
      'Exterior1st', 'Exterior2nd', 'ExterQual', 'ExterCond', 'Foundation',
      'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating',
      'HeatingQC', 'CentralAir', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
      'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
      'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
      'Functional', 'Fireplaces', 'GarageCars', 'GarageArea', 'PavedDrive',
      'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
      'ScreenPorch', 'PoolArea', 'MiscVal', 'MoSold', 'YrSold', 'SaleType',
      'SaleCondition', 'SalePrice'],
      dtype='object')
```

▼ EDA

```
df.nunique()
```

```
Id          1460
MSSubClass   15
MSZoning      5
LotArea    1073
Street        2
...
MoSold       12
YrSold        5
SaleType      9
SaleCondition  6
SalePrice    663
Length: 62, dtype: int64
```

```
df.Street.value_counts()
```

```
Pave    1454
Grvl      6
Name: Street, dtype: int64
```

```
df.MSZoning.value_counts()
```

```
RL      1151
RM       218
FV       65
RH       16
C (all)   10
Name: MSZoning, dtype: int64
```

```
df.YrSold.value_counts()
```

```
2009    338
2007    329
2006    314
2008    304
2010    175
Name: YrSold, dtype: int64
```

```
df.MoSold.value_counts()
```

```
6     253
7     234
5     204
4     141
8     122
3     106
10     89
11     79
9      63
12     59
1      58
2      52
Name: MoSold, dtype: int64
```

```
df = df.select_dtypes(exclude=['object'])
df
```

```
Id MSSubClass LotArea OverallQual OverallCond YearBuilt YearRemodAd
```

```
X = df.drop('SalePrice',axis = 1).values
Y = df.SalePrice.values
Y = np.reshape (Y, (-1,1))
```

```
print(X.shape)
print(Y.shape)
```

```
(1460, 34)
(1460, 1)
```

▼ Lets split the dataset into Training et Testing set

```
# We will split the data by using train_test_split from scikit
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.25,random_state=0)
```

```
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1095, 34)
(365, 34)
(1095, 1)
(365, 1)
```

▼ Feature Engineering

```
X = np.append(arr = np.ones((1460,1)), values = X, axis=1)
```

```
import statsmodels.api as sm
X_opt = X[:,[0,1,2,3,4,5]]
regressor = sm.OLS(endog = Y, exog = X_opt).fit()
regressor.summary()
```

OLS Regression Results

Dep. Variable: y **R-squared:** 0.666
Model: OLS **Adj. R-squared:** 0.665
Method: Least Squares **F-statistic:** 579.5
Date: Thu, 26 May 2022 **Prob (F-statistic):** 0.00
Time: 21:47:03 **Log-Likelihood:** -17744.
No. Observations: 1460 **AIC:** 3.550e+04
Df Residuals: 1454 **BIC:** 3.553e+04
Df Model: 5

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-9.22e+04	9029.698	-10.211	0.000	-1.1e+05	-7.45e+04
x1	1.2840	2.859	0.449	0.653	-4.325	6.893
x2	-162.5297	28.833	-5.637	0.000	-219.089	-105.970
x3	1.3523	0.123	11.024	0.000	1.112	1.593
x4	4.452e+04	880.507	50.561	0.000	4.28e+04	4.62e+04
x5	-775.8548	1088.662	-0.713	0.476	-2911.371	1359.661

Omnibus: 593.097 **Durbin-Watson:** 1.968

Prob(Omnibus): 0.000 **Jarque-Bera (JB):** 7698.739

Skew: 1.525 **Prob(JB):** 0.00

Kurtosis: 13.829 **Cond. No.** 1.09e+05

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.09e+05. This might indicate that there are strong multicollinearity or other numerical problems.

